




MICROSOFT BUILD






LAB520

Get Started with Models in Microsoft Foundry to Build AI Apps

A 5-Minute Introduction to Your Lab — Microsoft Build 2026



Meet Zava: Your Lab Scenario

You are Serena, a developer at of a global home-improvement retailer with thousands of daily product reviews.

Your mission: build an automated review moderation system that classifies customer reviews before they go live on the site.



Serena
AI Engineer/Developer

What is Microsoft Foundry?



Microsoft Foundry is your one-stop destination for discovering, evaluating, and deploying AI models.

Over 11k models from Microsoft, OpenAI, Meta, Mistral, Hugging Face, and more in one catalog.

Build, test, and ship production AI apps with identity-based auth, observability, and Responsible AI baked in.

Three concepts to know

Project

- Your workspace inside Foundry
- Scopes deployments, connections, and agents
- One project per environment is a good default

Deployment

- A model made callable in your project
- Has a name, region, and capacity
- You call the deployment — not the raw model

Hosted agent

- A packaged workflow (prompt + tools + model)
- Callable as an API, managed by Foundry
- The bridge from experiment to production

Your Lab Journey

-
- 1. Discover models in the Microsoft Foundry catalog
 - 2. Connect and send your first inference from Python
 - 3. Build a product review moderation app for Zava
 - 4. Deploy your workflow as a hosted agent

Getting Started with Skillable

- Your lab runs on Skillable a pre-configured virtual machine with everything installed
- Log into the VM using the credentials shown in the Skillable lab instructions
- Azure credentials, endpoints, and passwords are provided no personal subscription needed
- Open VS Code and Edge from the VM desktop and trust the workspace when prompted in VS Code.
- The lab repo is already cloned at C:\Users\LabUser\Desktop\Build26-LAB520-main

Your Skillable Lab Environment

Skillable provides all credentials in the lab instructions panel. Key items to locate:

Key Details

Auto Type – Keyboard Icon or Type will insert text directly into the VM.

VM Password – shown at the top of the lab instructions

Azure Username and Password – for signing into ai.azure.com

Project Endpoint – auto-configured in your .env file

Model Deployment Name – pre-deployed gpt-4.1-mini

What you'll walk away with

A working Python review moderation app powered by a Microsoft Foundry hosted model

Hands-on experience with the Foundry model catalog, inference API, and structured prompt design

A hosted agent deployed to your Foundry project — callable from any app, any runtime

Where to go next

Keep building

- Lab repo:
<http://github.com/microsoft/Build26-LAB520>
- Foundry samples:
<https://github.com/microsoft-foundry/foundry-samples>
- Foundry docs:
<https://learn.microsoft.com/azure/ai-foundry>
- Model catalog: <https://ai.azure.com>

Join the community

- Foundry Discord:
<https://aka.ms/foundry/discord>
- Foundry DevBlogs:
<https://devblogs.microsoft.com/foundry>
- Microsoft Build sessions:
<https://build.microsoft.com>
- Share what you built with
#MSFoundry #MSBuild



On the Build page:
Please give us feedback



<https://aka.ms/build2026-feedback>

THANK YOU
FOR YOUR FEEDBACK!

LAB520

Get Started with Models in Microsoft Foundry to Build AI Apps

Your feedback is very useful for us and
Helps us improve our labs for future events

Lab content: Please
open an issue in the repo



<https://aka.ms/build2026-LAB520>

We're with you every step of the way



Start building with Microsoft Foundry

ai.azure.com



Get the code from this session

aka.ms/build26-LAB520



Join our Discord community

aka.ms/foundry/discord